

A Four-Part Series

Knowing and Doing

On knowledge, markets, and the philosophy of practical competence

I. The Quarterback and the Spectator · II. Joe Economist Meets the Market · III. What the Market Actually Is · IV. The Marxist
Utopia Nobody Talks About

The Quarterback and the Spectator

On the difference between knowing that something is true and knowing how to do it

Every football fan has had the experience. The quarterback drops back, hesitates, throws an incompletion, and immediately — from the couch, from the upper deck, from the barstool — comes the verdict: he should have thrown to the tight end. It was obvious. Anyone could see it. And the fan, in his confidence, implicitly believes that he, given the relevant physical capacities, would have made the right call.

He is wrong. And the nature of his error runs considerably deeper than he realizes.

The most obvious factor is the informational asymmetry. The fan sees the entire field from above, or via television's carefully curated camera angles. The quarterback sees a collapsing pocket, bodies in motion at close range, and a narrow cone of field that shifts with every step. The fan's visual information and the quarterback's visual information are not the same information. But this, while real, is actually the most superficial element of the asymmetry.

The time problem

The quarterback operates inside a decision window of approximately two to three seconds. He does not deliberate. He cannot deliberate. What looks to the fan like a decision — a considered choice among alternatives — is in reality an act of highly trained pattern recognition executing under severe time constraint. The fan, watching the replay at leisure, is performing a leisured act of explicit reasoning and attributing its output to a domain where explicit reasoning is impossible. He is solving a different problem and claiming credit for solving the quarterback's problem.

This gap in temporal structure is deeper than it first appears. The psychologist Daniel Kahneman distinguishes between fast, automatic cognition and slow, deliberative cognition. The fan's judgment is a product of the slow system operating without pressure. The quarterback's execution must be a product of the fast system, which means that the relevant competence — the

competence actually on display in the game — is a competence the fan has never exercised and cannot evaluate by exercising its opposite.

The body is not optional

There is another factor the fan systematically ignores: the quarterback is simultaneously managing his body. Footwork, weight transfer, arm slot, grip pressure — these are not peripheral tasks that bracket the central cognitive problem. They compete for the same attentional bandwidth that the fan is entirely free to devote to reading the defense. The fan's body is at rest. He has no somatic noise. He is performing a purely cognitive task and imagining that the quarterback is performing the same task in the same way, only faster and under more pressure. But the somatic load does not just add difficulty — it qualitatively transforms the cognitive task.

Additionally, the quarterback is performing under conditions of genuine physiological arousal — elevated cortisol, narrowed perceptual field, altered working memory, compressed risk tolerance. These are not stylistic differences from the fan's neutral state. They are a different mode of cognition entirely. The fan cannot mentally transpose himself into that state by an act of imagination. This is what the behavioral economist George Loewenstein calls the empathy gap: our inability, from a cold state, to accurately simulate what we would do in a hot one.

The fan is solving a different, easier problem — and then attributing his solution to the harder one.

The philosopher's diagnosis

The English philosopher Gilbert Ryle, in his 1949 work *The Concept of Mind*, drew a distinction that cuts directly to the heart of this error. Ryle distinguished between knowing that and knowing how. Knowing that is propositional: it can be stated, written down, communicated. Knowing how is practical: it is a competence, an embodied skill, a capacity to execute that exists only in its exercise. You can know that a bicycle is steered by leaning into the turn before you know how to ride one. These are categorically different kinds of knowledge, and the possession of one confers no automatic claim on the other.

The fan has propositional knowledge of the situation. He can describe it, analyze it, evaluate it. He knows that the tight end was open, that the coverage indicated a blitz, that the correct read was the crossing route. He mistakes this for the practical, embodied competence that actual execution requires. He thinks that because he knows that, he would have known how.

This is a category error. And it is not a small one. Every other factor in the fan's misjudgment — the time pressure, the somatic load, the arousal state, the hindsight bias that makes the correct choice seem obvious after the fact, the counterfactual blindness that lets him imagine the tight end would have been open on the road not taken — every one of these is simply a specific dimension along which knowing that and knowing how come apart in this particular case.

The fan's error, stated in its most compressed form: he confuses his ability to evaluate performance with an ability to perform. These are not the same ability, and in most domains of genuine expertise, they are not even closely related.

NOTE

This distinction has consequences that extend far beyond football. In the essays that follow, we trace the same error through economics, financial markets, and political economy — where the stakes are considerably higher and the error considerably more costly.

Joe Economist Meets the Market

Why theoretical mastery of a domain is no preparation for operating within it

Consider Joe Economist. He has a PhD from a serious university. His dissertation concerned price dynamics in commodity markets. He has read everything — Samuelson, Fama, Shiller, Thaler. He understands equilibrium theory, behavioral finance, the efficient markets hypothesis and its critics. He is, by any reasonable measure, deeply knowledgeable about financial markets.

He decides to trade.

He loses money. Not immediately — he may even win for a while, which is worse, because it deepens his conviction that his knowledge is operative. But over time, over a sufficient number of trades, he underperforms. Often badly. The question is why.

The wrong kind of wrong

The naive answer is that his theories are incorrect. But this is not quite right, or at least not right in the way one might assume. Joe is not wrong about narrow matters of economic fact. He knows what interest rates are. He can read a balance sheet. His errors are not factual errors about specific data points. They are structural errors about the nature of the object he is operating on.

Joe's models were built on a set of idealizing assumptions that seemed innocuous in the academic context where they were developed but turn out to be catastrophic in the practical context where he is now deploying them. The most important of these is the assumption that price is a signal about value — that the current price of a security reflects the best available estimate of its fundamental worth, and that deviations from fundamental value are temporary aberrations that will correct.

This assumption is the intellectual heir of Graham and Buffett. It is, in a narrow sense, correct: a company's fundamental characteristics do constrain the range within which its shares will trade. But the range is wide. Very wide. And within that range, the assumption is not just imprecise but actively misleading — because it implies that price fluctuations carry information about value,

when in fact they primarily carry information about the current emotional state of the shareholder population.

Joe's models have the wrong probability architecture entirely. He has an accurate map of a fictional country.

The Gaussian catastrophe

Joe's models also embed, usually without examination, the assumption that returns are normally distributed — that large moves are rare and small moves are frequent, in the proportions described by the bell curve. This assumption is false. Financial returns exhibit fat tails: extreme events occur far more frequently than a Gaussian distribution predicts. The models are not merely imprecise about the tails. They are constitutively blind to them. When the tail event arrives, Joe's model has no response — it literally cannot see the event that is happening.

Nassim Taleb has made this point at length and with considerable acidity. It is worth stating plainly: the normal distribution assumption is not a simplification that can be patched. It is embedded in the mathematical foundations of most financial modeling. Removing it requires rebuilding from scratch, which academic economics has largely declined to do.

The circuit breaker problem

But the deepest problem with Joe's approach is not his distributional assumptions or his ontology of price. It is a practical problem that flows from both: his framework cannot generate an exit signal.

Joe enters a position because his model says the security is undervalued. The security declines. His model now says it is more undervalued than before — a better opportunity. He holds, or adds. The security declines further. The model reaches the same conclusion with greater conviction. There is no price level, short of catastrophic fundamental impairment, at which his model tells him he was wrong. Every adverse move is reinterpreted internally as a deeper discount. The framework has no circuit breaker because the framework treats every decline as confirmation rather than contradiction.

The sophisticated trader, by contrast, has a model coarse enough that reality can falsify it. The circuit breaker is not a separate rule bolted onto the model — it is internal to the model's structure, because the model was never overcommitted to a specific destination. Joe's model is so elaborate that it can absorb almost any adverse outcome without ever reaching the conclusion that it is wrong. This is not a feature of sophisticated modeling. It is the characteristic pathology of models that have been optimized for theoretical elegance rather than practical falsifiability.

Peru, 2002

The contrast is usefully illustrated by a historical case. In the early 2000s, Peru's Finance Minister was Pedro Pablo Kuczynski — a serious economist who had come up through the World Bank and Wall Street and who notably lacked a doctoral degree. The President, Alejandro Toledo, had a PhD in economics from Stanford.

Could Toledo have done Kuczynski's job? Almost certainly not — and the differential cannot be attributed to subject-matter ignorance, since both men were economists. What Kuczynski had that Toledo lacked was knowing-how in the specific domain that mattered: actual capital allocation, actual creditor negotiation, actual reading of market sentiment under pressure, in environments where being wrong had immediate and costly consequences. His formation had been in the world, not in representations of the world.

Toledo's Stanford PhD trained him to produce and evaluate theoretical representations of economies. This is knowing-that at its highest academic level. It is impressive and rigorous. It is also largely orthogonal to what a Finance Minister actually needs to do when foreign investors are spooked and the bond market is moving against you. Kuczynski's lack of a doctorate was not a deficit. It was, in a perverse sense, a credential — evidence that his competence had been formed in the field rather than in models of the field.

THE PATTERN IN BRIEF

Knowing-that: propositional, storable, communicable, transferable through instruction. Knowing-how: practical, embodied, formed through consequence-laden repetition, not transferable through instruction alone.

The quarterback fan has knowing-that about football. Toledo has knowing-that about economics. Joe Economist has knowing-that about markets. None of them has the knowing-how that the domain actually demands. The error is identical in structure across all three cases.

ESSAY THREE

What the Market Actually Is

Price as social fact, and why the blip-watchers are the object of study

The previous essay argued that Joe Economist fails in markets because his theoretical framework misrepresents the nature of the object he is operating on. This essay makes that claim precise. If price is not primarily a signal about value, what is it? And if fundamental analysis does not generate profitable predictions, what does?

Fundamentals determine a range, not a point

The first corrective to Joe's ontology is this: the fundamental characteristics of a company — its revenues, margins, competitive position, balance sheet — do not determine a correct price. They determine a range of legitimate prices. The range is wide. Within it, all prices are equally correct, because within it there is no fact of the matter about what the right price is. A blue-chip company does not have a price it should be trading at. It has a region of the price cosmos it inhabits, defined broadly by its fundamental quality, and within that region its actual price at any given moment is determined by something else entirely.

This has a corollary that Joe finds difficult to accept: since price fluctuations within the fundamental range carry no information about value, they carry no information that should update one's assessment of the investment. A decline does not mean the company is worth less. It means the current shareholders are in a different emotional state than they were yesterday. That is all it means.

Price is a social fact

The positive characterization follows directly. Price in a heavily traded security is a social fact — it is what the shareholder population collectively feels right now, expressed numerically. It is not a measurement of something that exists independently of that collective feeling. It is the feeling itself, quantified. This makes it more like a poll result than like a thermometer reading. A thermometer is wrong if it fails to track the temperature. A poll result cannot be wrong — it

simply is the distribution of opinion at that moment.

This reframing is not merely philosophical. It has immediate practical consequences. If price is a measurement, then adverse price movements should update your beliefs about the value of the underlying. If price is a social fact about emotional state, then adverse price movements are information about the current mood of the shareholder population — information that may, if you understand how moods evolve, be predictive of future price movements in a way that has nothing to do with fundamental value.

Joe thinks price is a measurement. The practical trader knows price is a social fact — what the blip-watchers collectively feel, expressed numerically.

The news means nothing

A closely related insight, which can be verified statistically and which experienced traders consistently rediscover: in heavily traded securities, news items do not reliably predict the direction of subsequent price moves. A negative analyst report is no more likely to push the price down than up. An inflation report that comes in hot may trigger a sell-off or may be shrugged off entirely depending on factors that have nothing to do with the content of the report.

This is not a paradox. It follows from the social-fact ontology. What moves prices is not the content of news but the reaction of the shareholder population to news, and that reaction depends on the current emotional state of the population, its current positioning, its current narrative, and a dozen other variables that the news item itself does not determine. The news is a trigger, not a cause. The same trigger produces radically different responses depending on the state of the system it enters.

The financial media, of course, works in precisely the opposite direction: it reverse-engineers a narrative from whatever price move has just occurred and presents it as the cause. Markets fell today because of concerns about interest rates. Markets rose today because of optimism about earnings. These narratives are largely confabulation. The move happened; a story was found to fit it.

The Goldman analyst problem

Joe Economist may concede some of this but retain a fallback position: even if retail investors are emotional and irrational, there are sophisticated institutional actors — Goldman, Bridgewater, the major hedge funds — whose activity keeps prices roughly tethered to fundamental value. These are the adults in the room. Their analysis is rigorous and their capital is large enough to matter.

This fallback is also wrong, and wrong in a specific way. The Goldman analyst covering semiconductor companies is a finance professional who has learned enough semiconductor vocabulary to write a credible report. The engineer who actually understands what a new process node means for competitive dynamics is in a design lab somewhere with no trading authority and no P&L.; The knowledge and the capital are structurally separated. The people who understand the technology cannot deploy capital on it. The people who can deploy capital do not understand the technology.

What Goldman is actually smart about — and it is genuinely smart about this — is flow, positioning, momentum, sentiment, and risk management. It is smart about the blips on the screen and the population of actors who move them. It is not smart about semiconductor physics. This means that Goldman's activity does not, in fact, keep prices tethered to fundamental technological value. It keeps prices tethered to what a large, sophisticated, well-informed community of blip-watchers collectively believes about what other blip-watchers will do next. That is a different thing.

Why heavy trading makes a security more predictable, not less

This leads to a counterintuitive insight: heavily traded securities are, in the relevant sense, more predictable than lightly traded ones. The standard assumption is the reverse — that a thinly traded small-cap, where a knowledgeable analyst might have genuine informational edge, is where predictability lives. But this gets the structure of the problem backwards.

A thinly traded security has uncertain fundamentals, sparse shareholder population, and idiosyncratic behavior that may be driven by one or two large holders or by private information that is unavailable to the public market. It is hard to model because the relevant variables are not observable.

A heavily traded security has a large, diversified, emotionally volatile shareholder population whose behavioral patterns are highly consistent. They panic-sell predictably. They chase green candles predictably. They are spooked by headline risk and reassured by narrative recovery in

ways that have proven stable across years and market cycles. The population is the object of study, and it is a well-characterized population. You are not trying to predict what the technology will do. You are predicting what a large group of human beings will do — and human beings, in aggregate, are remarkably consistent in their irrationalities.

THE PRACTICAL UPSHOT

Mean reversion on a heavily traded, structurally sound asset in an established sector is not a crude approximation of what sophisticated investors do. It is, stripped of pretension, what they actually do. The sophistication lies not in the model's complexity but in the precision of its self-knowledge: it knows exactly one thing — directional bias in a proven instrument — and declines to pretend it knows anything else.

ESSAY FOUR

The Marxist Utopia Nobody Talks About

On corporate socialism, para-statal entities, and the fiction of the private company

The previous essays argued that Joe Economist fails in markets because he misunderstands what price is and how it behaves. This essay argues that his failure runs still deeper: he misunderstands what large companies are. And this misunderstanding has consequences not only for how one trades but for how one thinks about markets, politics, and the economy as a whole.

Too big to fail is not an exception — it is the rule

The phrase "too big to fail" entered the public vocabulary during the 2008 financial crisis to describe the decision to rescue large banks rather than allow market discipline to operate. It was understood as an exceptional intervention — a departure from normal market functioning justified by extraordinary circumstances. This framing is wrong. The crisis made visible a subsidy that had always existed. It was not a new policy. It was the revelation of a permanent condition.

Large companies — not only banks, but any company that has achieved genuine systemic scale — are continuously, invisibly backstopped by the network of dependencies they have helped create. Suppliers extend favorable terms because losing the customer would be devastating to their own operations. Creditors provide cheap capital because the alternative is a disruption they cannot afford. Regulators move slowly and carefully because aggressive action would have costs that extend far beyond the company being regulated. Governments intervene at the limit because the counterfactual is worse for everyone.

None of this appears on any balance sheet. All of it supports the share price. It is a permanent, diffuse, structural subsidy that exists independently of the company's quarterly earnings. It is not charity. It is the rational self-interest of every actor in a network that needs the node to remain functional.

The public/private distinction dissolves at scale

Mainstream economics depends on a clean separation between public actors (governments) and private actors (companies). The entire apparatus of market theory — price signals, competition, efficient allocation, the role of profit as a coordinating mechanism — is built on this distinction. Private actors pursue private interests in a domain cleanly separable from state power. Government sets the rules; companies play the game.

This picture is a reasonable approximation for small and medium enterprises. For very large companies it is not an approximation. It is a fiction that both parties — government and corporation — have strong interests in maintaining.

The large company staffs the regulatory agencies that oversee it, through the revolving door between corporate leadership and government positions. It writes, through lobbying and direct participation in the legislative process, the regulations that govern its competitive environment. It provides infrastructure — communications, logistics, financial — that the state depends on to function. It is, in practice, a co-author of the rules of the game it is supposedly playing. The notion that it stands outside the state, subject to it, disciplined by it, is not a description of reality. It is a story both parties find convenient to tell.

The large company gets the best of both worlds: state protection and private gain. It is a third thing that economic theory has no clean category for.

The para-statal entity

A more accurate description than either "public" or "private" is para-statal. The large company is a branch of governance that has been granted a special dispensation: operate like a government agency — with all the structural protection, implicit guarantees, and systemic indispensability that entails — except when it suits you to operate as a private profit-maximizing entity. The dispensation is not written down anywhere. It does not need to be. It emerges from the mutual dependencies that large scale creates.

The consequences extend to domains that seem remote from economics. A sufficiently large platform company that controls speech infrastructure is not meaningfully a "private" actor when it makes content moderation decisions. Its decisions have the functional character of state action — they determine what speech is permitted in the spaces where public discourse actually occurs

— while retaining the legal character of private corporate policy. The constitutional protections that citizens have against government censorship do not apply, because the entity exercising censorship is technically private. This is not a quirk or an edge case. It is the predictable consequence of allowing communications infrastructure to be organized as for-profit enterprise at systemic scale.

The Marxist utopia

Socialism, stripped to its operational core, means the socialization of risk: the insulation of certain actors from the consequences of failure, with losses absorbed by the collective. It means resource allocation determined partly by political position rather than purely by market performance. It means guaranteed survival for actors deemed essential to the system's functioning.

Large companies in advanced capitalist economies have all of this. Their losses are socialized — explicitly in crises, implicitly at all times through regulatory capture, subsidized capital, and network support. Their survival is partially guaranteed by their structural indispensability. Their competitive position is maintained partly through political relationships. They do not fully face the market. The market faces them.

What they have achieved is, in certain respects, more robust than anything a political socialist program has ever delivered. The USSR could not guarantee Leningrad's survival the way the Federal Reserve implicitly guarantees JPMorgan's. The large company's utopia — socialized risk, privatized gain, permanent structural protection — is more durably engineered than any political utopia has managed to be.

The small company, meanwhile, lives in something closer to the rawer version of capitalism that textbooks describe — genuinely exposed, genuinely disciplined by market forces, genuinely on its own when it fails. Nobody socializes its losses. Nobody captures regulators on its behalf. It faces something closer to the theoretical free market than any large company has faced in a century.

The investment implication

This analysis has a direct and under-appreciated implication for how one thinks about large-cap equity. The share price of a systemically embedded company is not simply a claim on its future earnings. It is a claim on its future earnings plus the shadow value of its permanent structural

protection plus the option value of its illegible asset base plus the implicit guarantee of its continued existence by a network of actors who need it to continue existing. None of these latter components appears in any discounted cash flow model. All of them support the floor.

Joe Economist, modeling large companies as autonomous profit-maximizing entities operating in competitive markets, is missing most of what is actually determining their valuations. He is not wrong about the part he can see. He is blind to the larger structure that makes the part he can see possible.

THE DEEPEST IRONY

The political actors who most loudly champion free markets and oppose socialism are almost universally the representatives of large capital — the entities most thoroughly insulated from market discipline. And the political actors who most loudly champion socialism are often trying to extend to ordinary people a fraction of the protection that large companies already permanently enjoy. The actual existing socialism in Western economies is corporate socialism. It is never named as such because the naming would be politically catastrophic for everyone involved.

What classical economics calls a large private company is a category error. The thing it is pointing at — a genuinely private actor operating purely on market signals — essentially does not exist above a certain size threshold. Below that threshold you have real private companies. Above it you have para-statal entities that perform privateness when convenient. Trading as if this were not true is trading with a map that has the wrong legend.
